
Introduction to Message Passing Interface (MPI)

Message Passing Interface

- API for distributed memory programming
 - Now de-fact standard in HPC/parallel computing
 - Explicitly write data communication
 - Message Passing
 - What is a message:
 - “data” and additional information
- We will use MPI to parallelize codes in our exercise problems

MPI APIs

- ❑ Initialize MPI environment
- ❑ Query process information
- ❑ 1 to 1 communication
 - Synchronous/Asynchronous
 - Blocking/Non-blocking
- ❑ Collective communication
 - Broadcast
 - Reduction
 - Synchronization

Initialize MPI

- ❑ Must call following two functions

```
#include "mpi.h"

int main(int argc, char *argv[] )
{
    MPI_Init(&argc, &argv);

    // application body

    MPI_Finalize();
    return 0;
}
```

Query process information

- Each process has own ID number

```
#include "mpi.h"

int main(int argc, char *argv[] )
{
    int np, id;
    MPI_Init(&argc, &argv);

    // Query the number of processes
    MPI_Comm_size(MPI_COMM_WORLD, &np);
    // Query own ID
    MPI_Comm_rank(MPI_COMM_WORLD, &id);

    MPI_Finalize();
    return 0;
}
```

Used APIs

- ❑ `MPI_Init(int *, char **)`
 - Initialize MPI environment
- ❑ `MPI_Finalize()`
 - Finish MPI environment
- ❑ `MPI_COMM_WORLD (MPI_Comm struct)`
 - It is called a communicator
 - We use the communicator to access process information
 - This is a default communicator
- ❑ `MPI_Comm_size(MPI_Comm, int *)`
 - Query the number of processes in a given communicator
- ❑ `MPI_Comm_rank(MPI_Comm, int *)`
 - Query own ID

Hello World example

□ Compile

- `mpicc hello.c -o hello`

□ Execution

- `mpirun -np ### executable`

□ Example with 4 processes

```
std4dc1[~] mpicc hello.c -o hello
std4dc1[~] mpirun -np 4 ./hello
Hello world from process 2 (4)
Hello world from process 1 (4)
Hello world from process 0 (4)
Hello world from process 3 (4)
std4dc1[~]
```

Single Program Multiple Data

- ❑ MPI is a type of SPMD parallel programming
 - Source code is a single file.
 - An Execution file is also single.
 - ❑ The Program working on multiple data
 - We launch multiple processes running the same program in different memory space.
 - We need “data communication” between processes

SPMD example (1)

```
int np, id;
...
MPI_Comm_size(MPI_COMM_WORLD, &np);
MPI_Comm_rank(MPI_COMM_WORLD, &id);
...

switch(id) {
    case 0:
        do something assigned id = 0
        break
    case 1:
        do something assigned id = 1
        break
    case 2:
        do something assigned id = 2
        break
}
```

SPMD example (2)

□ Ex. 1

- Calculation of pi

$$\int_0^1 \frac{4 \, dx}{1 + x^2}$$

```
h = 1.0/(double)n;  
sum = 0.0;  
  
// main loop  
for (i = 0; i < n; i++) {  
    x = h*((double)i - 0.5);  
    sum += f(x);  
}  
mypi = h*sum;
```

- We have “n” evaluations of f(x).
 - We have summation of the partial sum.
- How do we parallelize this?

SPMD example (3)

- We divide the loop into “np” parts.
 - Serial: $i = 0$ to $n-1$
 - Parallel: $i = s[id]$ to $e[id]$

```
// main loop
for (i = s[id]; i < e[id]; i++) {
    x = h*((double)i - 0.5);
    sum += f(x);
}
```

- $n = 1000$, $np = 2$
 - $s[0] = 0$, $e[0] = 500$
 - $s[1] = 500$ $e[1] = 1000$
- $n = 1000$, $np = 3$
 - $s[0] = 0$, $e[0] = 333$
 - $s[1] = 333$ $e[1] = 666$
 - $s[2] = 666$ $e[1] = 1000$ (treatment of last index)

SPMD example (4)


□ A SPMD way solution

```
int s[128], e[128];
```

```
...
```

```
s[id] = (n/np)*id;
```

```
e[id] = (n/np)*(id+1) + (id-1 == np ? n % np : 0);
```



□ We can calculate each partial sum in parallel

□ How about the total sum?


■ We use communication!!

Measure elapsed time in MPI

- ❑ `double MPI_Wtime(void)`
 - Time measurement function in MPI.
 - Return value is the elapsed in seconds.

```
double time_start, time_end;
...
time_start = MPI_Wtime();
for (i = s[id]; i < e[id]; i++) {
    x = h*((double)i - 0.5);
    sum += f(x);
}
time_end = MPI_Wtime();
// time_end-time_start is the elapsed time for the loop
```

Summation




- ❑ Processes $\text{id} = 1$ to $n-1$ send own partial sum to process $\text{id}=0$ (main process)
- ❑ The main process calculate the total sum

```
if (id != 0) {  
    send own partial sum to the main process  
} else {  
    receive partial sums from other processes  
    calculate the total sum  
}
```

1 to 1 Communication


- ❑ We use ID in communication
- ❑ In 1 to 1 communication, we need to match send and receive operations
 - Sender :sends data to process (id=0)
 - Receiver :receives data from process (id=1)




Collective communications

□ Collective communications

■ Broadcast




■ Reduction



Synchronization


□ Synchronization



1 to 1 Communication details(1)

- Send : MPI_Send()
- Receive : MPI_Recv()

```
MPI_Status status;  
int data;  
  
if (id == 0) {  
    data = 3;  
    MPI_Send(&data, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);  
}  
  
if (id == 1) {  
    MPI_Recv(&data, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);  
}
```



1 to 1 Communication details(2)

► MPI_Send:

```
MPI_Send(void *data, int count, MPI_Datatype type,  
          int dest, int tag, MPI_Comm comm)
```

C/C++

- **data:** memory location containing data
- **count:** number of elements of type “type” to send
- **type:** MPI type of each element
- **dest:** destination rank for the data
- **tag:** identification of message
- **comm:** communicator

1 to 1 Communication details(3)

► MPI_Recv:

```
MPI_Recv(void *data, int count, MPI_Datatype type,  
          int source, int tag, MPI_Comm comm, MPI_Status * status)
```

C/C++

- **data:** memory location intended for data
- **count:** number of elements of type “type” to send
- **type:** MPI type of each element
- **source:** rank to receive data from
Special rank: **MPI_ANY_SOURCE**
- **tag:** identification of message to receive
Special tag: **MPI_ANY_TAG**
- **comm:** communicator
- **status:** status of the return
Special status: **MPI_STATUS_IGNORE**